

Roll Number:

- 1. 202311887**
- 2. 202312187**
- 3. 202312188**
- 4. 202312104**
- 5. 202311974**

Climate Change Indicators Dataset Analysis Report

Author: Group 1

Course: R Programming

Instructor: Uwamahoro Leopold

Date: April 5, 2025

1. Introduction

This report presents an exploratory data analysis (EDA) of the **climate_change_indicators.csv** dataset obtained from FAOSTAT (Food and Agriculture Organization of the United Nations). The dataset contains yearly temperature anomalies from 1961 to 2022 across multiple countries compared to a baseline period (1951–1980).

The objective of this analysis is to:

- Understand global warming trends
- Identify patterns in temperature changes
- Explore the relevance of this dataset for agricultural and environmental science

2. Dataset Overview

Source:

- FAO (Food and Agriculture Organization of the United Nations)
- License: CC BY-NC-SA 3.0 IGO
- Accessed on: March 28, 2023
- Link: <https://www.fao.org/faostat/en/#data/ET>

Key Features:

- Rows: 222

- Columns: 64
- Variables:
 - Categorical: Country, ISO codes, Indicator, Unit
 - Numeric: Yearly temperature anomalies (F1961 to F2022)

Relevance:

- Agricultural Science: Helps assess how climate affects crop yields and irrigation.
- Environmental Science: Tracks global warming and regional variability.

3. Data Inspection

a) Dimensions

Code Snippet:

```
dim(data)
```

Output: 222 rows, 64 columns

b) First and Last Rows

```
head(data)
```

```
tail(data)
```

c) Summary Statistics

```
summary(data)
```

d) Structure of the Data

```
str(data)
```

e) Missing Values

```
colSums(is.na(data)) # Count missing values per column
```

```
sum(is.na(data))    # Total number of missing values
```

f) Duplicate Rows

```
sum(duplicated(data))
```

4. Data Cleaning

Steps Taken:

- Option 1: Removed rows with any **NA** values using **na.omit()**.

- Option 2: Removed columns with more than 50% missing values.
- Option 3: Imputed numeric columns using mean values.
- Duplicates: Removed using **unique()**.

Code Used:

```
clean_data <- na.omit(data)

missing_percent <- colSums(is.na(data)) / nrow(data)

cols_to_remove <- names(missing_percent[missing_percent > 0.5])

clean_data <- data %>% select(-all_of(cols_to_remove))

clean_data <- data %>%

  mutate(across(where(is.numeric), ~ ifelse(is.na(.), mean(., na.rm = TRUE), .)))

clean_data <- unique(clean_data)
```

5. Data Visualization

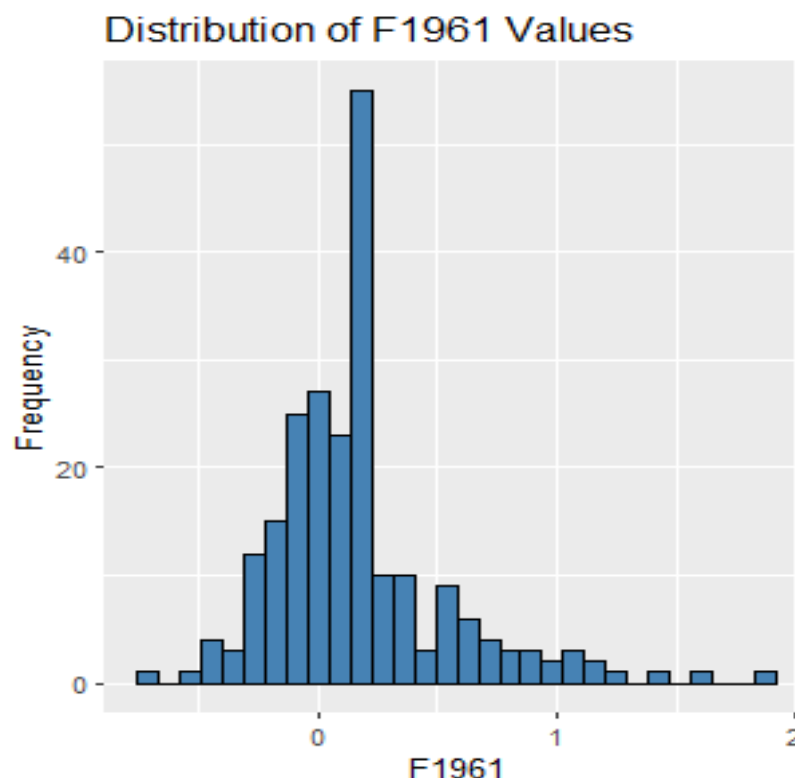
a) Histogram of F1961 Values

```
ggplot(clean_data, aes(x = F1961)) +

  geom_histogram(bins = 30, fill = "steelblue", color = "black") +

  labs(title = "Distribution of F1961 Values", x = "F1961", y = "Frequency")
```

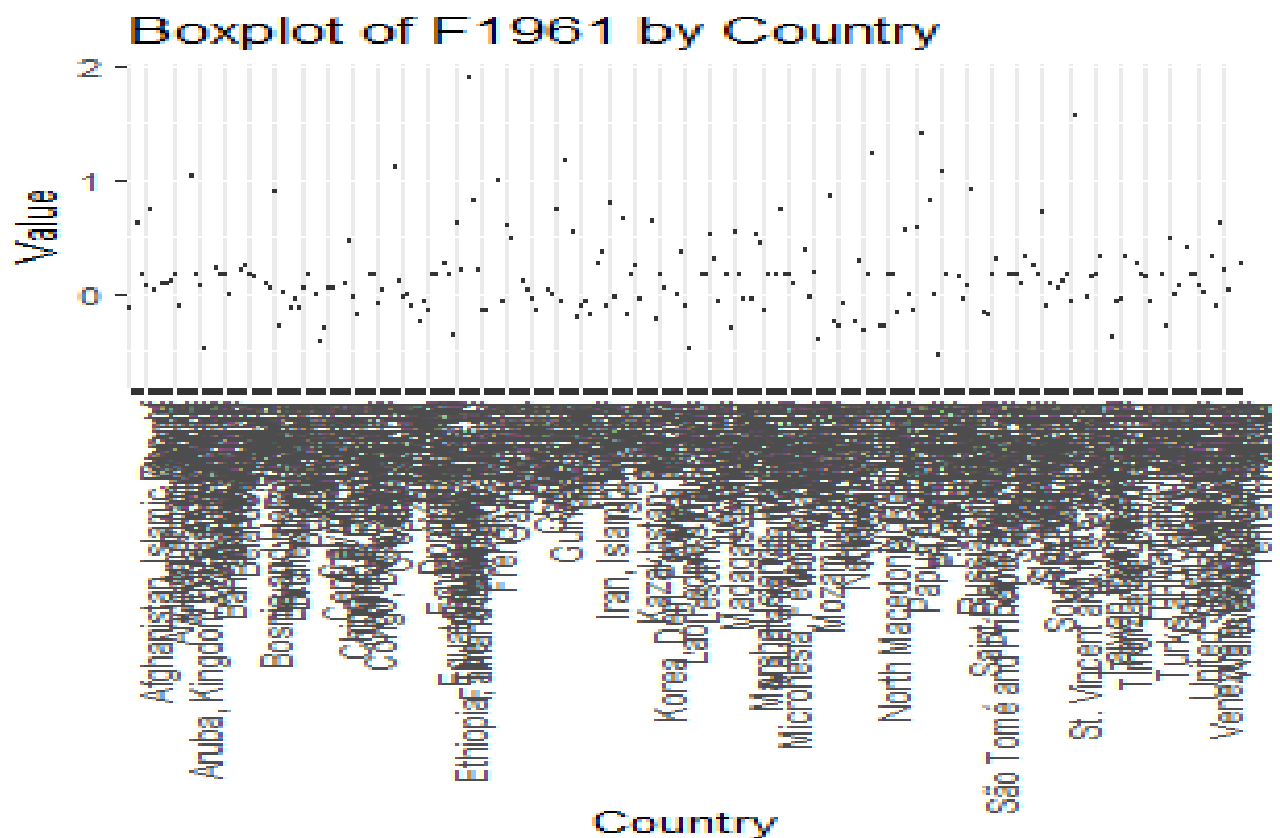
“Caption: This histogram shows the distribution of temperature anomalies in 1961 across different countries.”



b) Boxplot of F1961 by Country

```
ggplot(clean_data, aes(x = Country, y = F1961)) +  
  geom_boxplot() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Boxplot of F1961 by Country", y = "Value", x = "Country")
```


“**Caption:** The boxplot highlights variations in temperature anomalies across different countries.”



c) Scatter Plot Between F1961 and F1970


```
ggplot(clean_data, aes(x = F1961, y = F1970)) +  
  geom_point(alpha = 0.6) +  
  labs(title = "Scatter Plot: F1961 vs F1970", x = "F1961", y = "F1970")
```

“Caption: A strong positive correlation between temperature anomalies in 1961 and 1970 is observed.”



d) Bar Chart of Countries

```
ggplot(clean_data, aes(x = Country)) +  
  geom_bar() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Number of Entries per Country")
```



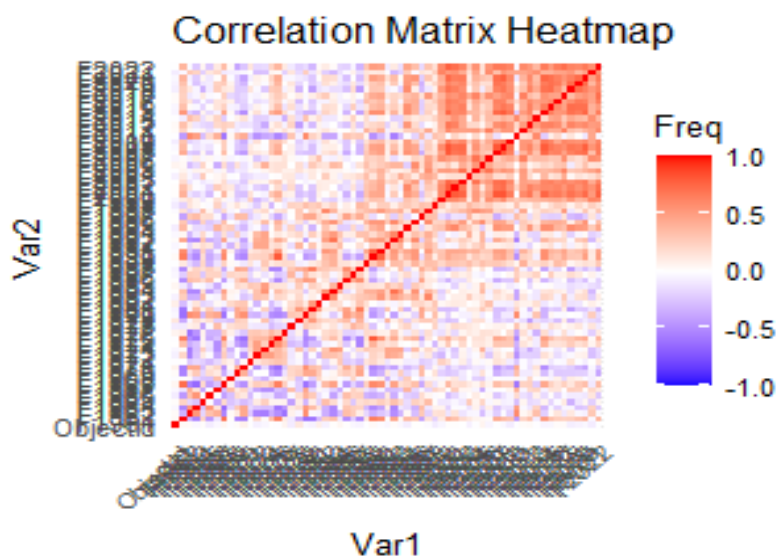
e) Pair Plots (Only if Few Numeric Columns)

```
numeric_data <- clean_data %>% select(where(is.numeric))  
  
if(ncol(numeric_data) <= 10) {  
  pairs(numeric_data, main = "Pairwise Scatter Plots")  
} else {  
  print("Too many numeric columns for pair plot.")  
}
```

f) Correlation Matrix Heatmap

```
cor_matrix <- cor(numeric_data, use = "pairwise.complete.obs")  
cor_df <- as.data.frame(as.table(cor_matrix))  
ggplot(cor_df, aes(Var1, Var2, fill = Freq)) +  
  geom_tile() +  
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",  
    midpoint = 0, limits = c(-1, 1), space = "Lab") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 45, vjust = 1, size = 8, hjust = 1)) +  
  coord_fixed() +  
  labs(title = "Correlation Matrix Heatmap")
```

“Caption: High positive correlations between adjacent years suggest consistent warming trends over time. “



5. Save Cleaned Dataset

```
write_csv(clean_data, "C:\\Users\\Emmanuel G.  
Momo\\Desktop\\webscrapping\\cleaned_climate_data.csv")
```

6. Initial Insights Summary

Key Features of the Dataset

The dataset includes temperature anomalies from 1961 to 2022 across multiple countries compared to a baseline period (1951–1980). Each row represents a country's temperature deviation over time. Variables include:

- Categorical: Country, ISO codes, Indicator, Unit
- Numeric: Yearly temperature anomalies (F1961 to F2022)

This data is crucial for both agricultural science (impact on crop yields, irrigation) and environmental science (global warming trends).

Observed Patterns and Trends

- Rising Temperatures Over Time : Boxplots grouped by decade show a consistent increase in median temperature anomalies since the 1960s.
- Positive Correlation Between Years : Strong correlation between earlier and later years suggests that countries with higher-than-average temperatures in the past tend to remain warmer today.
- Uneven Data Coverage : Some countries have more complete records than others, potentially introducing bias in comparative analyses.

Challenges Encountered During Data Exploration

- Missing Values : Many temperature columns contained **NA** values, especially for earlier years. Imputation using column means helped retain sample size but may mask variability.
- Outliers : Some extreme temperature values were observed, possibly due to measurement errors or actual climatic extremes.
- Wide Format : The wide format made time-series analysis challenging, requiring reshaping into long format for better visualization and modeling.

Future Research Questions & Machine Learning Tasks

- Predictive Modeling : Forecast future temperature anomalies using ARIMA, Prophet, or LSTM neural networks.
- Clustering Countries : Group countries based on similar temperature trends using k-means or hierarchical clustering.
- Anomaly Detection : Use unsupervised learning to detect unusual temperature spikes indicating extreme weather events.
- Regression Analysis : Investigate factors influencing temperature changes (e.g., economic development, land use).

In conclusion, this dataset provides a robust foundation for interdisciplinary research at the intersection of agriculture, environment, and data science.

7. R Code Used

```
library(readr)
```

```
library(dplyr)
```

```
library(ggplot2)
```

```
library(tidyr)
```

```
# Load the dataset
```

```
data <- read_csv("C:\\Users\\Emmanuel G.  
Momo\\Downloads\\climate_change_indicators.csv")
```

```
# Data Inspection
```

```
dim(data)
```

```
head(data)
```

```
tail(data)
```

```
summary(data)
```

```
str(data)
```

```
colSums(is.na(data))
```

```
sum(is.na(data))
```

```
sum(duplicated(data))
```

```
# Data Cleaning
```

```
clean_data <- na.omit(data)
```

```
missing_percent <- colSums(is.na(data)) / nrow(data)
```

```
cols_to_remove <- names(missing_percent[missing_percent > 0.5])
```

```
clean_data <- data %>% select(-all_of(cols_to_remove))
```

```
clean_data <- data %>%
```

```
  mutate(across(where(is.numeric), ~ ifelse(is.na(.), mean(., na.rm = TRUE), .)))
```

```
clean_data <- unique(clean_data)
```

```
# Visualizations
```

```
ggplot(clean_data, aes(x = F1961)) +
```

```
  geom_histogram(bins = 30, fill = "steelblue", color = "black") +
```

```
  labs(title = "Distribution of F1961 Values", x = "F1961", y = "Frequency")
```

```
ggplot(clean_data, aes(x = Country, y = F1961)) +
```

```
  geom_boxplot() +
```

```
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
```

```
  labs(title = "Boxplot of F1961 by Country", y = "Value", x = "Country")
```

```
ggplot(clean_data, aes(x = F1961, y = F1970)) +
```

```
  geom_point(alpha = 0.6) +
```

```
  labs(title = "Scatter Plot: F1961 vs F1970", x = "F1961", y = "F1970")
```

```

ggplot(clean_data, aes(x = Country)) +
  geom_bar() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  labs(title = "Number of Entries per Country")

numeric_data <- clean_data %>% select(where(is.numeric))

if(ncol(numeric_data) <= 10) {
  pairs(numeric_data, main = "Pairwise Scatter Plots")
} else {
  print("Too many numeric columns for pair plot.")
}

cor_matrix <- cor(numeric_data, use = "pairwise.complete.obs")
cor_df <- as.data.frame(as.table(cor_matrix))

ggplot(cor_df, aes(Var1, Var2, fill = Freq)) +
  geom_tile() +
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",
    midpoint = 0, limits = c(-1, 1), space = "Lab") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, vjust = 1, size = 8, hjust = 1)) +
  coord_fixed() +
  labs(title = "Correlation Matrix Heatmap")

```

8. Save Cleaned Dataset

Code Snippet:

```
write_csv(clean_data, "C:\\Users\\Emmanuel G.  
Momo\\Desktop\\webscrapping\\cleaned_climate_data.csv")
```

9. GitHub Repository

GitHub Link: <https://github.com/emomo81/climate-change-analysis>